



Day 3: Scheduling Algorithms (CSIS 2260 Preview)



Goals

- Understand the purpose of CPU scheduling in multitasking environments.
- Learn three major scheduling algorithms: FCFS, SJF, and Round Robin.
- Visualize how scheduling affects CPU efficiency and process wait times.
- Practice building Gantt charts and simulations.



Topics to Cover



Why Scheduling Matters

- The OS decides **which process gets CPU time** and **when**.
- A good scheduler:
 - Maximizes CPU utilization.
 - Minimizes waiting and turnaround time.
 - Ensures fairness among processes.



Key Scheduling Algorithms

1. First-Come, First-Served (FCFS)

- Processes are handled in the order they arrive.
- Simple, but can cause **convoy effect** (long waits behind big jobs).

2. Shortest Job First (SJF)

- Executes the process with the shortest burst time first.

- Optimal in terms of average waiting time, but hard to predict burst time.

3. Round Robin (RR)

- Each process gets a fixed time slice (quantum).
- Preemptive scheduling—great for time-sharing systems.

Watch (Optional, 15–20 min)

- [CPU Scheduling Basics](#)
- [Shortest Job First Scheduling.\(Solved Problem 1\)](#)

Do

1. Gantt Chart Practice

- Draw a Gantt chart for each algorithm using the following sample:

1. FCFS

Process	Arrival	Burst	Turnaround	Waiting
-----	-----	-----	-----	-----
P1	0	5	$5 - 0 = 5$	$5 - 5 = 0$
P2	1	3	$8 - 1 = 7$	$7 - 3 = 4$
P3	2	1	$9 - 2 = 7$	$7 - 1 = 6$
P4	3	2	$11 - 3 = 8$	$8 - 2 = 6$

Gantt Chart

0 P1 5 P2 8 P3 9 P4 11

Average waiting time = $(0 + 4 + 6 + 6) / 4 = 3.5$

2. SRTF - Shortest Remaining Time First, Preemptive SJF

- It's definitely worth understanding, even if not common in Windows/Linux kernels.

Process	Arrival	Burst	Turnaround	Waiting
P1	0	5	$11 - 0 = 11$	$11 - 5 = 6$
P2	1	3	$7 - 1 = 6$	$6 - 3 = 3$
P3	2	1	$3 - 2 = 1$	$1 - 1 = 0$
P4	3	2	$5 - 3 = 2$	$2 - 2 = 0$

Gantt Chart

0 P1 1 P2 2 P3 3 P4 5 P2 7 P1 11

Average waiting time = $(6 + 3 + 0 + 0) / 4 = 2.25$

3. SJF - Non-Preemptive

Process	Arrival	Burst	Turnaround	Waiting
P1	0	5	$5 - 0 = 5$	$5 - 5 = 0$
P2	1	3	$11 - 1 = 10$	$10 - 3 = 7$
P3	2	1	$6 - 2 = 4$	$4 - 1 = 3$
P4	3	2	$8 - 3 = 5$	$5 - 2 = 3$

Gantt Chart

0 P1 5 P3 6 P4 8 P2 11

Average waiting time = $(0 + 7 + 3 + 3) / 4 = 3.25$

4. RR

Process	Arrival	Burst	Turnaround	Waiting
P1	0	5	$11 - 0 = 11$	$11 - 5 = 6$
P2	1	3	$10 - 1 = 9$	$9 - 3 = 6$

P3		2		1		$5 - 2 = 3$		$3 - 1 = 2$
P4		3		2		$9 - 3 = 6$		$6 - 2 = 4$

Gantt Chart

0 P1 2 P2 4 P3 5 P1 7 P4 9 P2 10 P1 11

Average waiting time = $(6 + 6 + 2 + 4) / 4 = 4.5$

- Try FCFS, SJF (non-preemptive), and Round Robin (quantum = 2).
- Calculate:
 - Turnaround Time = Completion Time - Arrival Time
 - Waiting Time = Turnaround Time - Burst Time

2. Spreadsheet or Code Simulation (Optional)

- Use Excel, Google Sheets, or a Python/JavaScript snippet to simulate Round Robin logic.
- [Practice it with function](#)
- [Practice it with Class](#)

```
class Process:
    def __init__(self, pid, arrival_time, burst):
        self.pid = pid
        self.arrival_time = arrival_time
        self.burst = burst
        self.remaining_time = burst
        self.is_visited = False
        self.completion_time = 0
        self.turnaround_time = 0
        self.waiting_time = 0
```

```
class RoundRobinScheduler:
    def __init__(self, processes, quantum):
```

```

self.processes = processes
self.quantum = quantum
self.time = 0
self.queue = []
self.gantt_chart = []

def is_finished(self):
    return all(p.remaining_time == 0 for p in
self.processes)

def is_newly_arrived(self, p):
    return not p.is_visited and p.arrival_time <=
self.time

def enqueue_arrived_processes(self):
    for process in self.processes:
        if self.is_newly_arrived(process):
            self.queue.append(process)
            process.is_visited = True

def run(self) -> None:
    while not self.is_finished():
        # enqueue
        self.enqueue_arrived_processes()
        # wait
        if not self.queue:
            self.time += 1
            continue
        # pop first queue
        current = self.queue.pop(0)
        # CPU running
        self.gantt_chart.append((self.time, current.pid))
        # Update
        if current.remaining_time > self.quantum:
            self.time += self.quantum

```

```

        current.remaining_time -= self.quantum
    else:
        self.time += current.remaining_time
        current.remaining_time = 0
        current.completion_time = self.time
        # Refresh : enqueue current process after new arrival
        self.enqueue_arrived_processes()
        if current.remaining_time > 0:
            self.queue.append(current)
# wrap up
        self.gantt_chart.append((self.time, "Finished"))

```

Checkpoint

- I can describe how each scheduling algorithm works.
- I can draw a correct Gantt chart for FCFS, SJF, and RR.
- I calculated waiting and turnaround time for a process.
- I understand how different algorithms affect performance.